

Reads- from:

T_1 reads from T_2
(on item A)
if, in the schedule,
we have

$w_2(A), \dots, r_1(A)$
without any writes
in between



T_1 reads from T_2 ,
then T_2 gets
rolled back
— what do we
do with T_1 ,
then?



$w_2(A), \dots, r_1(A), \dots, \underline{c_1}, \underline{c_2}$

↑
reads from



Recoverable : Intuition :

⇒ potentially
"cascading
rollbacks" if T_2
needs to be rolled back



Another example of
recoverable \mathcal{S} :

$w_2(A), \dots, \underline{c_2}, \dots, r_1(A), \dots$

reads from c_1



Avoids cascading
rollbacks (ACR) ✓
schedule:

each T reads
from a committed
transaction



ACR: intuition:

$w_2(A), \dots, r_i(A), \dots, r_j(A)$



reads
from

c_2 must
be here



Not ACR:

reads-from

... $w_2(A), \dots, w_2(B), \dots, r_1(B), \dots$

reads-from

..., $c_2, \dots, r_1(A), \dots, c_1$



Relations

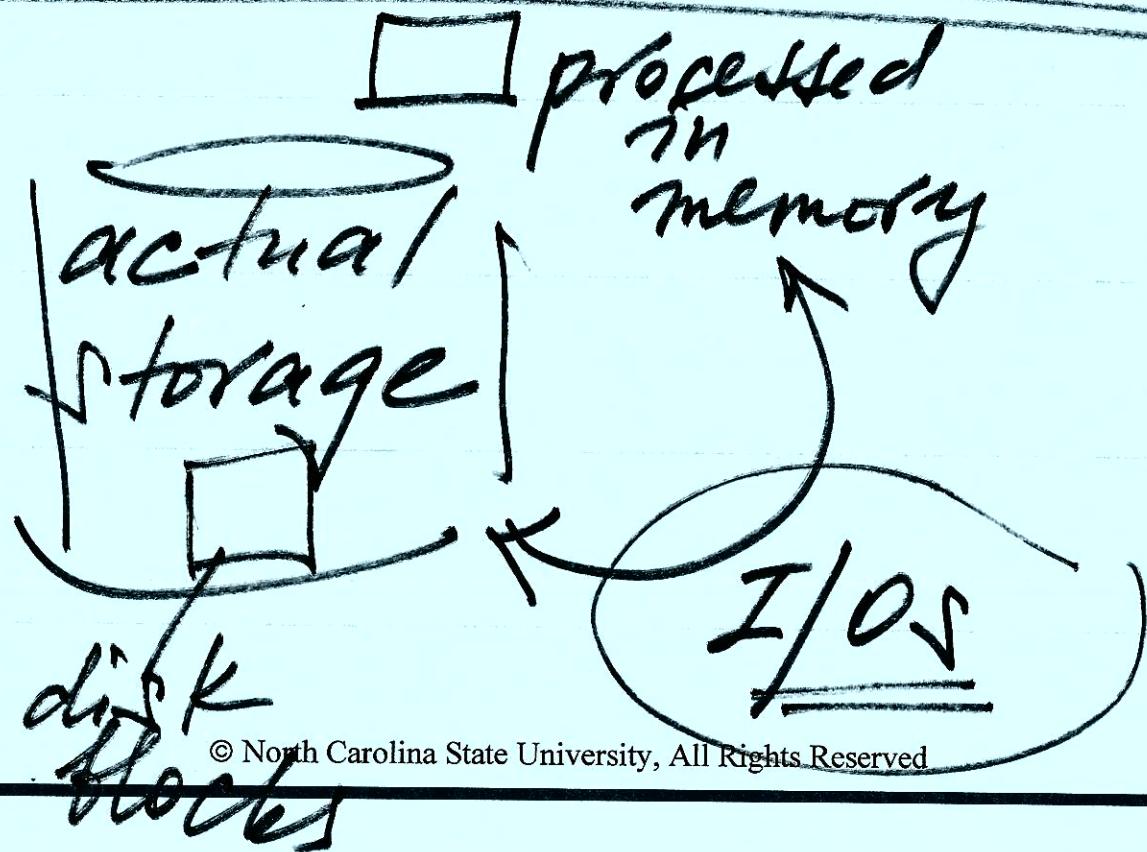


queries:

SQL

user side:

DBMS
side:



Summary:

Translate SQL into relational algebra
(because procedural),
then optimize the resulting expression
for your complexity measure



Intuition:

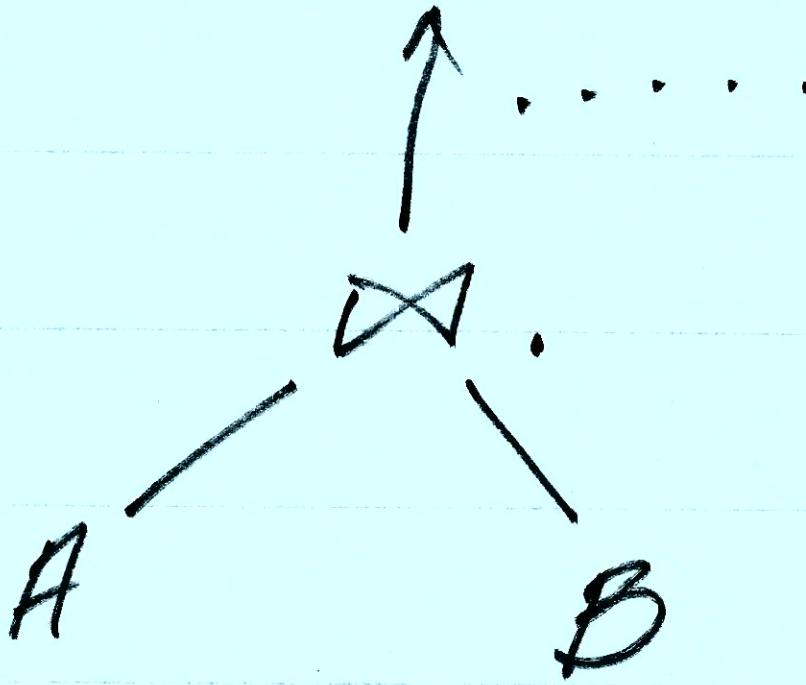
SQL Stmt

mapped \rightarrow default
rd-alg
expression

Select X, P, JV
From A, B
Where C

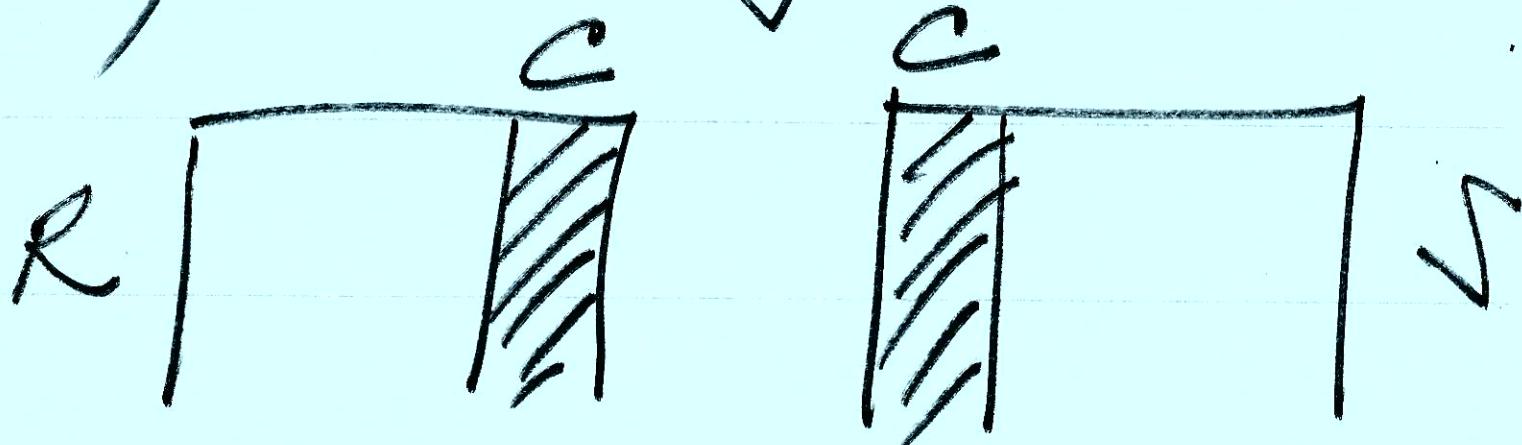
$\vdash_{TQ,B,TX} \Delta$





Natural join α

(1) Hash join: $R \bowtie S$

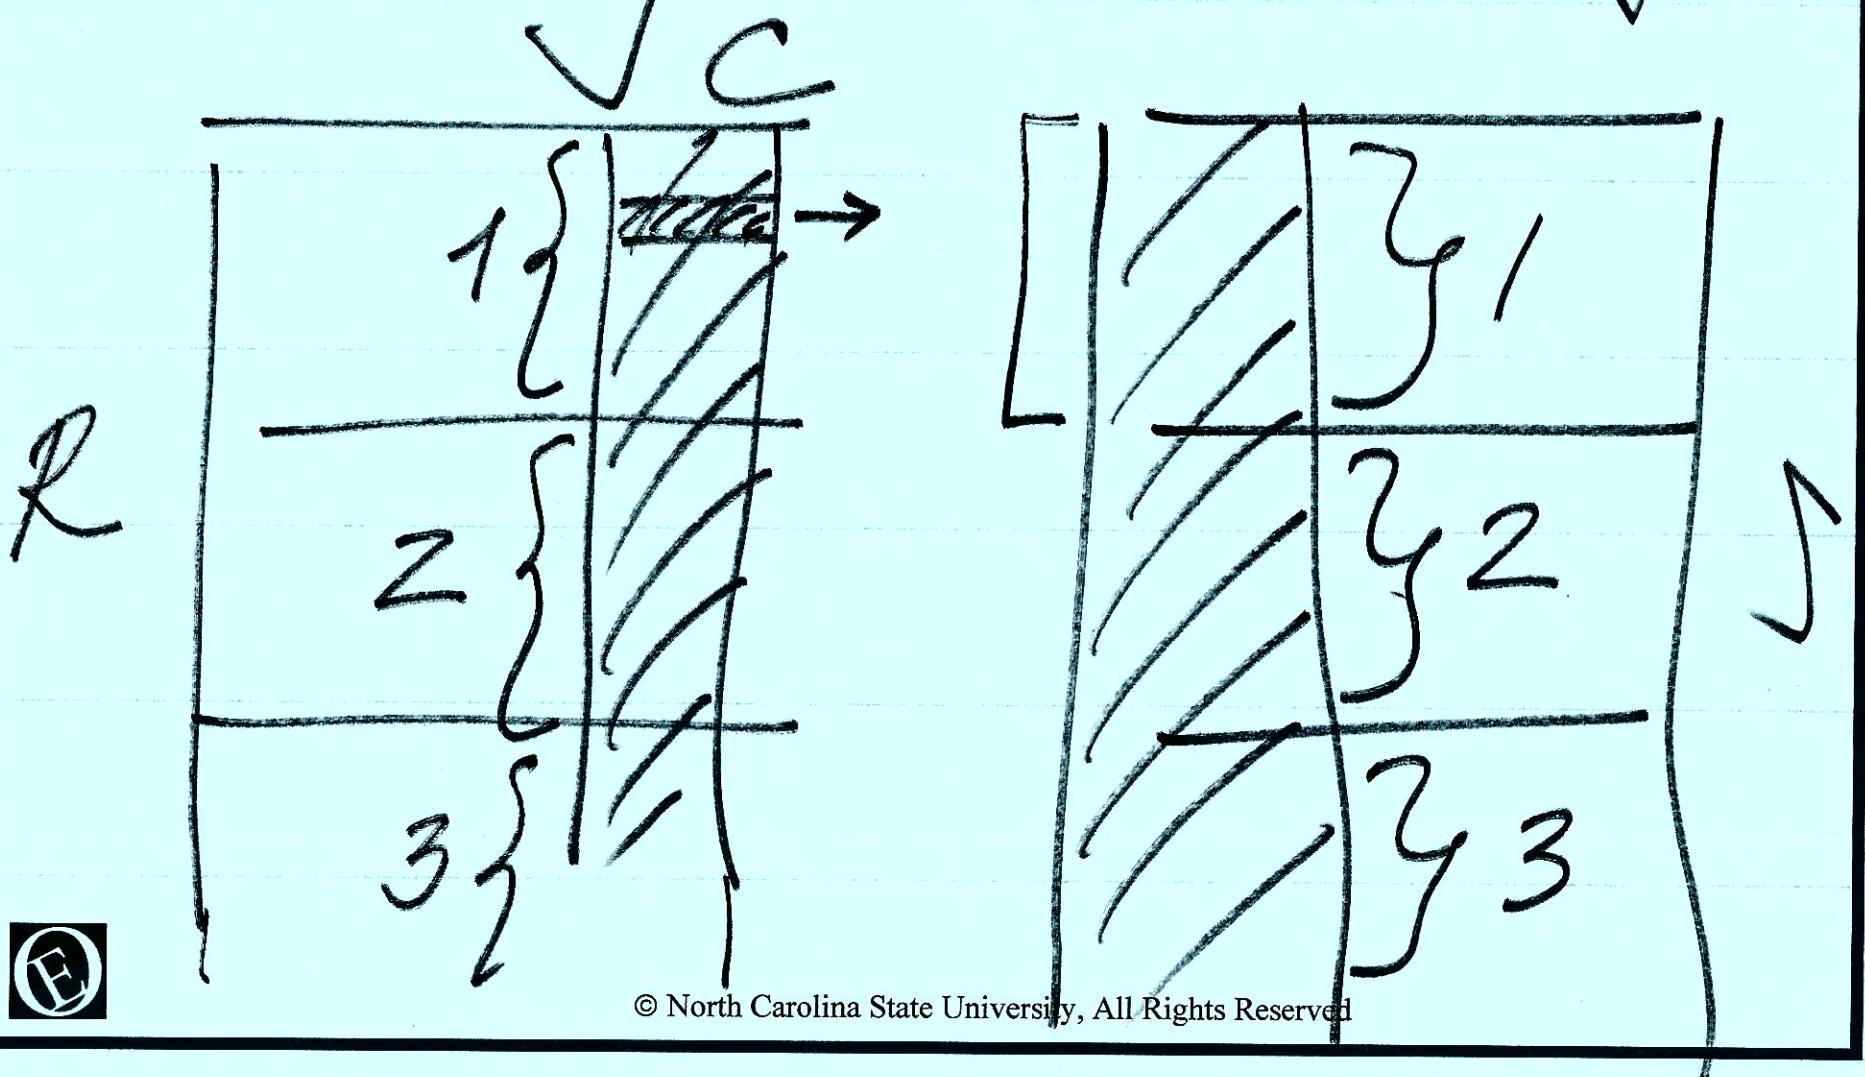


$c=1 | c=2 | c=3 | c=4$ | bucketize on
 c
 1 scan of R into memory

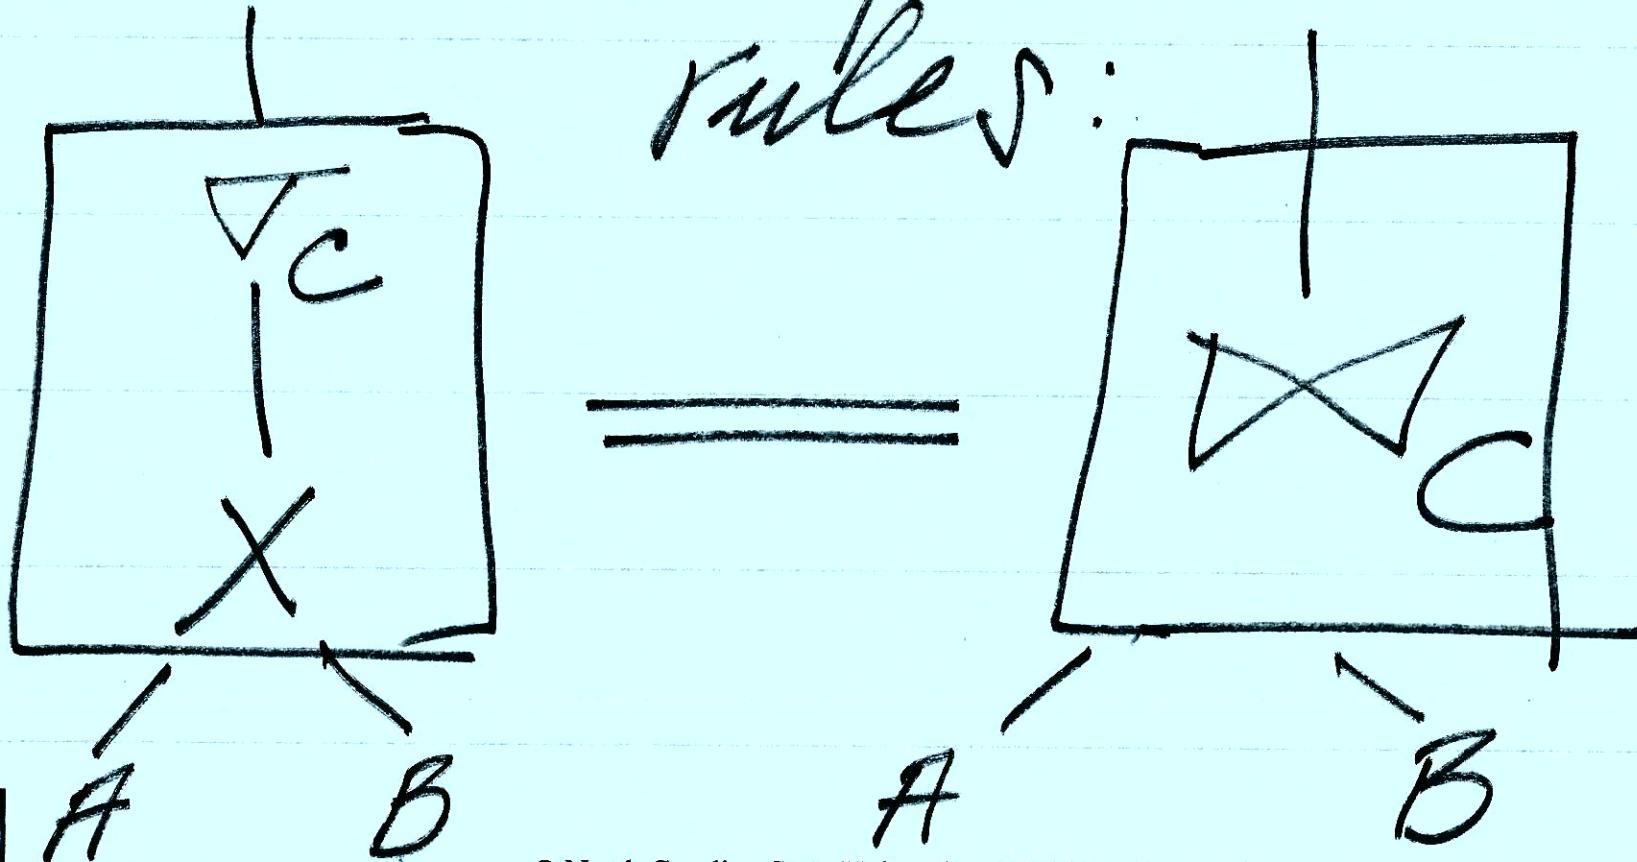


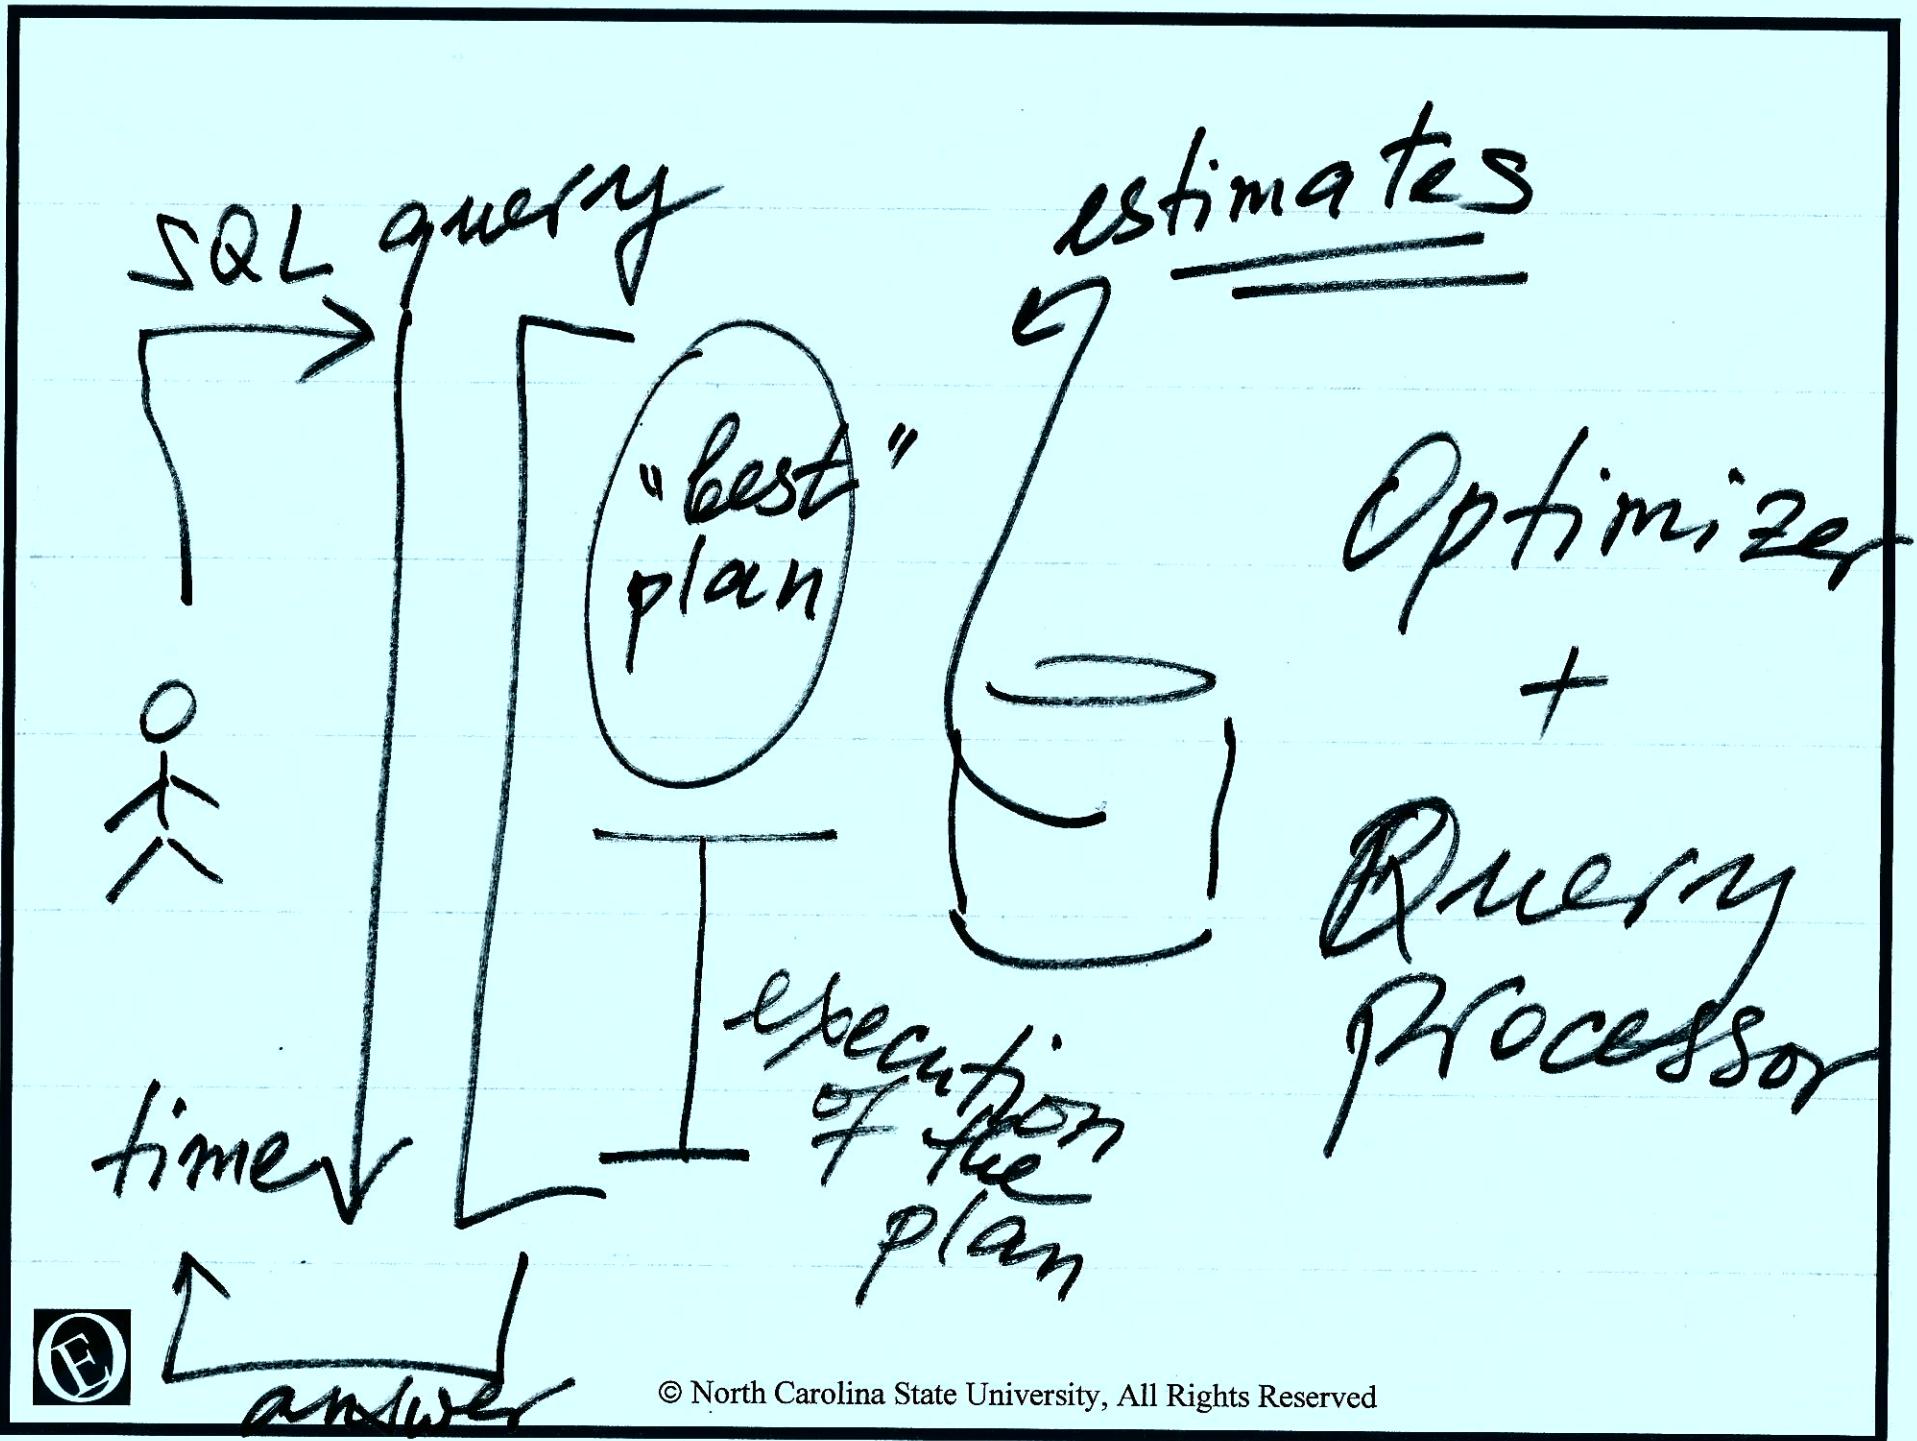
(2) Sort - merge

join: $R \bowtie S$



Optimizers:
Transformation
rules:





Keys to optimization:

- (1) work with only necessary data
- (2) process the data efficiently at each stage

